

```
In [23]: import pandas as pd
import numpy as np
import re
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
```

```
In [7]: data = pd.read_csv('/Users/raunavsharma/Downloads/IMDB.csv')
print(data.shape)
data.head(10)
```

(50000, 2)

Out[7]:

	review	sentiment
--	--------	-----------

0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production. The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive
5	Probably my all-time favorite movie, a story o...	positive
6	I sure would like to see a resurrection of a u...	positive
7	This show was an amazing, fresh & innovative i...	negative
8	Encouraged by the positive comments about this...	negative
9	If you like original gut wrenching laughter yo...	positive

```
In [8]: data.describe()
```

Out[8]:

	review	sentiment
count	50000	50000
unique	49582	2

	review	sentiment
top	Loved today's show!!! It was a variety and not...	negative
freq	5	25000

```
In [10]: data["sentiment"].value_counts()
```

```
Out[10]: negative    25000
         positive    25000
         Name: sentiment, dtype: int64
```

```
In [86]: reviews_train = []
         for line in open('/Users/raunavsharma/Downloads/movie_data/full_train.txt', 'r'):
             reviews_train.append(line.strip())

         reviews_test = []
         for line in open('/Users/raunavsharma/Downloads/movie_data/full_test.txt', 'r'):
             reviews_test.append(line.strip())
```

```
In [87]: replace_no_space = re.compile("[. ,:;!'\?,\\"()\\[\\]]")
         replace_with_space = re.compile("(<br\\s*/><br\\s*/>)|(\\-)|(\\/)")

         def preprocess_reviews(reviews):
             reviews = [replace_no_space.sub("", line.lower()) for line in reviews]
             reviews = [replace_with_space.sub(" ", line) for line in reviews]

             return reviews

         reviews_train_clean = preprocess_reviews(reviews_train)
         reviews_test_clean = preprocess_reviews(reviews_test)
```

```
In [88]: p = CountVectorizer(binary=True)
         p.fit(reviews_train_clean)
         X = p.transform(reviews_train_clean)
         X_test = p.transform(reviews_test_clean)
         X_train = p.transform(reviews_test_clean)
```

```
In [122]: target = [1 if i < 12500 else 0 for i in range(25000)]
```

```

X_train, X_val, y_train, y_val = train_test_split(
    X, target, train_size = 0.75
)

for c in [0.01, 0.05, 0.25, 0.5, 1]:

    lr = LogisticRegression(C=c)
    lr.fit(X_train, y_train)
    print ("Accuracy for C=%s: %s"
          % (c, accuracy_score(y_val, lr.predict(X_val))))

```

Accuracy for C=0.01: 0.86736

/Users/raunavsharma/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

Accuracy for C=0.05: 0.87744

/Users/raunavsharma/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

Accuracy for C=0.25: 0.8752

/Users/raunavsharma/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

Accuracy for C=0.5: 0.87312

Accuracy for C=1: 0.8696

```
/Users/raunavsharma/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
In [123... final_model = LogisticRegression(C=0.05)  
final_model.fit(X, target)  
print ("Final Accuracy: %s"  
      % accuracy_score(target, final_model.predict(X_test)))
```

Final Accuracy: 0.88156

```
/Users/raunavsharma/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
In [124... feature_to_coef = {  
    word: coef for word, coef in zip(  
        p.get_feature_names(), final_model.coef_[0]  
    )  
}  
for best_positive in sorted(  
    feature_to_coef.items(),  
    key=lambda x: x[1],  
    reverse=True)[:5]:  
    print (best_positive)  
  
for best_negative in sorted(  
    feature_to_coef.items(),  
    key=lambda x: x[1][:5]:  
    print (best_negative)
```

```
('excellent', 0.9287863620357822)
('perfect', 0.7916864122112646)
('great', 0.6740677270263481)
('amazing', 0.6131983883926218)
('superb', 0.6011322324407015)
('worst', -1.3645720549264835)
('waste', -1.166786439535322)
('awful', -1.0321101486702577)
('poorly', -0.8751951963515281)
('boring', -0.8567832150328295)
```

```
In [ ]: ## In this project, I first cleaned the data and then performed the Logistic Regression. I also got various accuracies
## References:
## https://towardsdatascience.com/building-a-logistic-regression-in-python-step-by-step-becd4d56c9c8
## https://www.datacamp.com/community/tutorials/understanding-logistic-regression-python
```